

In []:

```
!pip install tensorflow-addons
```

Requirement already satisfied: tensorflow-addons in /usr/local/lib/python3.7/dist-packages (0.17.0)
Requirement already satisfied: typeguard>=2.7 in /usr/local/lib/python3.7/dist-packages (from tensorflow-addons) (2.7.1)
Requirement already satisfied: packaging in /usr/local/lib/python3.7/dist-packages (from tensorflow-addons) (21.3)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from packaging->tensorflow-addons) (3.0.9)

In []:

```
#General imports
from __future__ import print_function
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import progressbar
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import normalize
from sklearn.decomposition import PCA

# Keras imports
import keras
from tensorflow.keras import layers

from keras.preprocessing.image import load_img
from keras.models import Sequential, Model
from tensorflow.keras.optimizers import *
from keras.utils.np_utils import to_categorical
import keras.backend as K

# application (model) imports
from tensorflow.keras import applications
#from keras.applications.inception_v3 import preprocess_input
from keras.layers import Dense

import tensorflow as tf
import tensorflow_datasets as tfds
import tensorflow_addons as tfa
```

In []:

```
# load the data from Tensorflow
tfds.disable_progress_bar()
(train, test), info = tfds.load(
    "stanford_dogs",
    split=["train", "test"],
    as_supervised=True, # Include labels
    with_info = True,
    try_gcs = True,
)

print("Number of training samples: %d" % tf.data.experimental.cardinality(train))
print("Number of test samples: %d" % tf.data.experimental.cardinality(test))
```

Number of training samples: 12000
Number of test samples: 8580

In []:

```
print(info)
```

```
tfds.core.DatasetInfo(
    name='stanford_dogs',
```

```

version=0.2.0,
description='The Stanford Dogs dataset contains images of 120 breeds of dogs from aro
und
the world. This dataset has been built using images and annotation from
ImageNet for the task of fine-grained image categorization. There are
20,580 images, out of which 12,000 are used for training and 8580 for
testing. Class labels and bounding box annotations are provided
for all the 12,000 images.',
homepage='http://vision.stanford.edu/aditya86/ImageNetDogs/main.html',
features=FeaturesDict({
    'image': Image(shape=(None, None, 3), dtype=tf.uint8),
    'image/filename': Text(shape=(), dtype=tf.string),
    'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=120),
    'objects': Sequence({
        'bbox': BBoxFeature(shape=(4,), dtype=tf.float32),
    }),
}),
total_num_examples=20580,
splits={
    'test': 8580,
    'train': 12000,
},
supervised_keys=('image', 'label'),
citation="""@inproceedings{KhoslaYaoJayadevaprakashFeiFei_FGVC2011,
author = "Aditya Khosla and Nityananda Jayadevaprakash and Bangpeng Yao and
Li Fei-Fei",
title = "Novel Dataset for Fine-Grained Image Categorization",
booktitle = "First Workshop on Fine-Grained Visual Categorization,
IEEE Conference on Computer Vision and Pattern Recognition",
year = "2011",
month = "June",
address = "Colorado Springs, CO",
}
@inproceedings{imagenet_cvpr09,
AUTHOR = {Deng, J. and Dong, W. and Socher, R. and Li, L.-J. and
Li, K. and Fei-Fei, L.},
TITLE = {{ImageNet: A Large-Scale Hierarchical Image Database}},
BOOKTITLE = {CVPR09},
YEAR = {2009},
BIBSOURCE = "http://www.image-net.org/papers/imagenet_cvpr09.bib"}""",
redistribution_info=,
)

```

In []:

```

plt.figure(figsize=(10, 10))
for i, (image, label) in enumerate(train.take(9)):
    ax = plt.subplot(3, 3, i + 1)
    plt.imshow(image)
    plt.title(int(label))
    plt.axis("off")

```

36



118



46




103



113



116





70



9



13



In []:

```
#resize the augmented images to 224 * 224

size = (224, 224)

train = train.map(lambda x, y: (tf.image.resize(x, size), y))
test = test.map(lambda x, y: (tf.image.resize(x, size), y))
#we don't need to resize the test image
```

In []:

```
size = (224, 224, 3)

train = train.map(lambda x, y: (tf.image.random_flip_left_right(x), y))
train = train.map(lambda x, y: (tf.image.random_crop(x, size), y))
#train = train.map(lambda x, y: (tf.image.random_saturation(x, size), y))
```

In []:

```
#batch the data and use caching & prefetching to optimize loading speed

batch_size = 64

train = train.batch(batch_size).prefetch(buffer_size=10)
test = test.batch(batch_size).prefetch(buffer_size=10)
```

In []:

```
#set up base model
input_shape = (224, 224, 3)
num_classes=120

base_model = applications.resnet_v2.ResNet101V2(input_shape=input_shape, include_top=False, weights='imagenet', pooling='avg')

x = base_model.output #We use Keras Functional API here
predictions = Dense(num_classes, activation='softmax')(x)
model = Model(inputs = base_model.input, outputs=predictions)
```

In []:

```
# compile the model
# Optimizer: Nesterov's Accelerated Gradient
alpha = 0.0001 #weight decay
momentum = 0
model.compile(optimizer=tfa.optimizers.SGDW(momentum = momentum, weight_decay = alpha, nesterov=True, learning_rate=0.01),
              loss=['sparse_categorical_crossentropy'], metrics=['sparse_categorical_accuracy'])
```

In []:

```
#train the model
```

```
history = model.fit(train, epochs=15, validation_data= test, batch_size=256)
```

Epoch 1/15

188/188 [=====] - 140s 613ms/step - loss: 2.9108 - sparse_categorical_accuracy: 0.3708 - val_loss: 1.3919 - val_sparse_categorical_accuracy: 0.5963

Epoch 2/15

188/188 [=====] - 131s 698ms/step - loss: 1.0210 - sparse_categorical_accuracy: 0.7369 - val_loss: 1.0962 - val_sparse_categorical_accuracy: 0.6922

Epoch 3/15

188/188 [=====] - 111s 590ms/step - loss: 0.5373 - sparse_categorical_accuracy: 0.8737 - val_loss: 1.0512 - val_sparse_categorical_accuracy: 0.6974

Epoch 4/15

188/188 [=====] - 111s 591ms/step - loss: 0.3173 - sparse_categorical_accuracy: 0.9373 - val_loss: 1.0014 - val_sparse_categorical_accuracy: 0.7101

Epoch 5/15

188/188 [=====] - 111s 592ms/step - loss: 0.1968 - sparse_categorical_accuracy: 0.9689 - val_loss: 1.0008 - val_sparse_categorical_accuracy: 0.7100

Epoch 6/15

188/188 [=====] - 111s 593ms/step - loss: 0.1318 - sparse_categorical_accuracy: 0.9837 - val_loss: 0.9681 - val_sparse_categorical_accuracy: 0.7210

Epoch 7/15

188/188 [=====] - 132s 700ms/step - loss: 0.0906 - sparse_categorical_accuracy: 0.9912 - val_loss: 0.9837 - val_sparse_categorical_accuracy: 0.7157

Epoch 8/15

188/188 [=====] - 112s 593ms/step - loss: 0.0646 - sparse_categorical_accuracy: 0.9956 - val_loss: 0.9954 - val_sparse_categorical_accuracy: 0.7189

Epoch 9/15

188/188 [=====] - 111s 592ms/step - loss: 0.0510 - sparse_categorical_accuracy: 0.9974 - val_loss: 0.9973 - val_sparse_categorical_accuracy: 0.7168

Epoch 10/15

188/188 [=====] - 131s 699ms/step - loss: 0.0410 - sparse_categorical_accuracy: 0.9987 - val_loss: 0.9950 - val_sparse_categorical_accuracy: 0.7224

Epoch 11/15

188/188 [=====] - 131s 699ms/step - loss: 0.0361 - sparse_categorical_accuracy: 0.9982 - val_loss: 1.0009 - val_sparse_categorical_accuracy: 0.7219

Epoch 12/15

188/188 [=====] - 111s 591ms/step - loss: 0.0328 - sparse_categorical_accuracy: 0.9985 - val_loss: 1.0111 - val_sparse_categorical_accuracy: 0.7203

Epoch 13/15

188/188 [=====] - 111s 593ms/step - loss: 0.0288 - sparse_categorical_accuracy: 0.9993 - val_loss: 1.0065 - val_sparse_categorical_accuracy: 0.7234

Epoch 14/15

188/188 [=====] - 111s 592ms/step - loss: 0.0273 - sparse_categorical_accuracy: 0.9994 - val_loss: 1.0090 - val_sparse_categorical_accuracy: 0.7227

Epoch 15/15

188/188 [=====] - 131s 699ms/step - loss: 0.0253 - sparse_categorical_accuracy: 0.9996 - val_loss: 1.0151 - val_sparse_categorical_accuracy: 0.7216

In []:

```
#Define plot function
```

```
def plot_loss_accuracy(history):
```

```
    historydf = pd.DataFrame(history.history, index=history.epoch)
```

```
    plt.figure(figsize=(8, 6))
```

```
    historydf.plot(ylim=(0, max(1, historydf.values.max())))
```

```
    loss = history.history['loss'][-1]
```

```
    acc = history.history['sparse_categorical_accuracy'][-1]
```

```
    val_acc = history.history['val_sparse_categorical_accuracy'][-1]
```

```
    val_error = (1 - val_acc)
```

```
    plt.title('Loss: %.3f, sparse_categorical_accuracy: %.3f' % (loss, acc))
```

```
    print('Validation Error: %.3f' % (val_error))
```

In []:

```
plot_loss_accuracy(history)
```

Validation Error: 0.278

<Figure size 576x432 with 0 Axes>

Loss: 0.025, sparse_categorical_accuracy: 1.000

